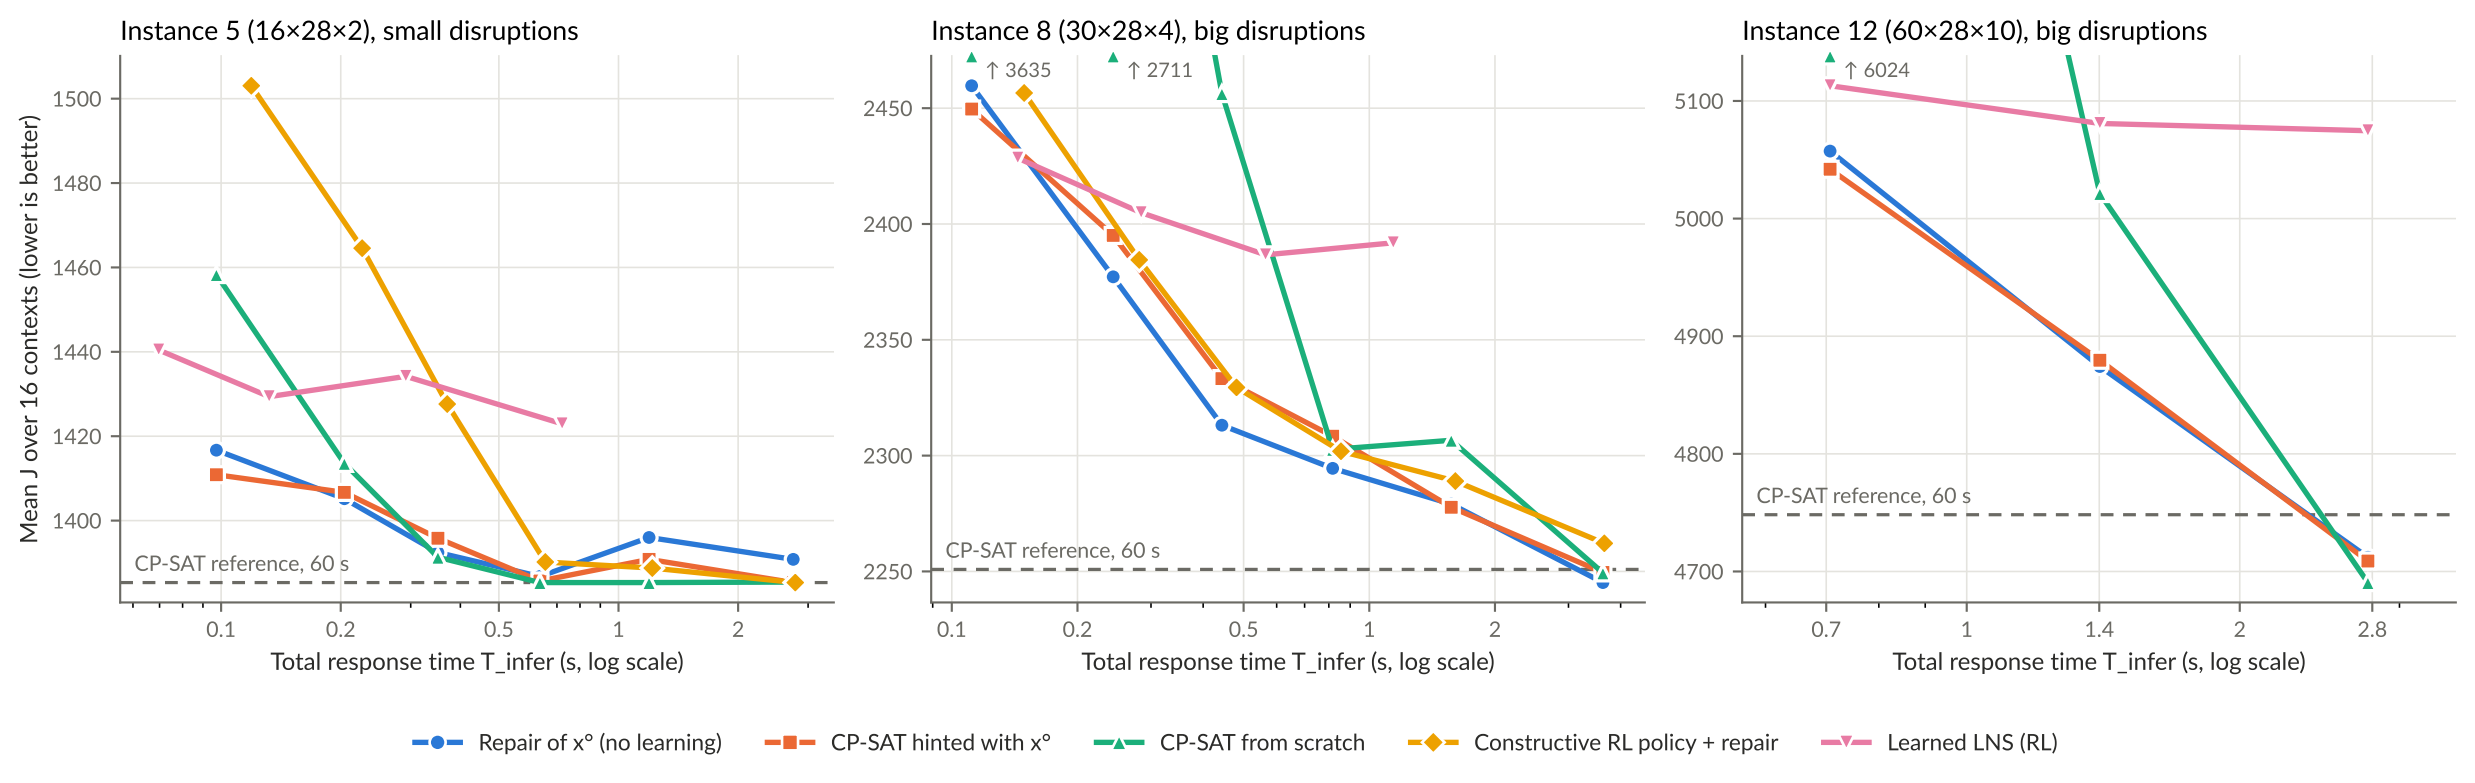
Instance 5 (16×28×2), small disruptions
Instance 8 (30×28×4), big disruptions
Instance 12 (60×28×10), big disruptions
Mean J over 16 contexts (lower is better)
Total response time T_infer (s, log scale)
CP-SAT reference, 60 s
↑ 3635
↑ 2711
↑ 6024
Repair of x° (no learning)
CP-SAT hinted with x°
CP-SAT from scratch
Constructive RL policy + repair
Learned LNS (RL)