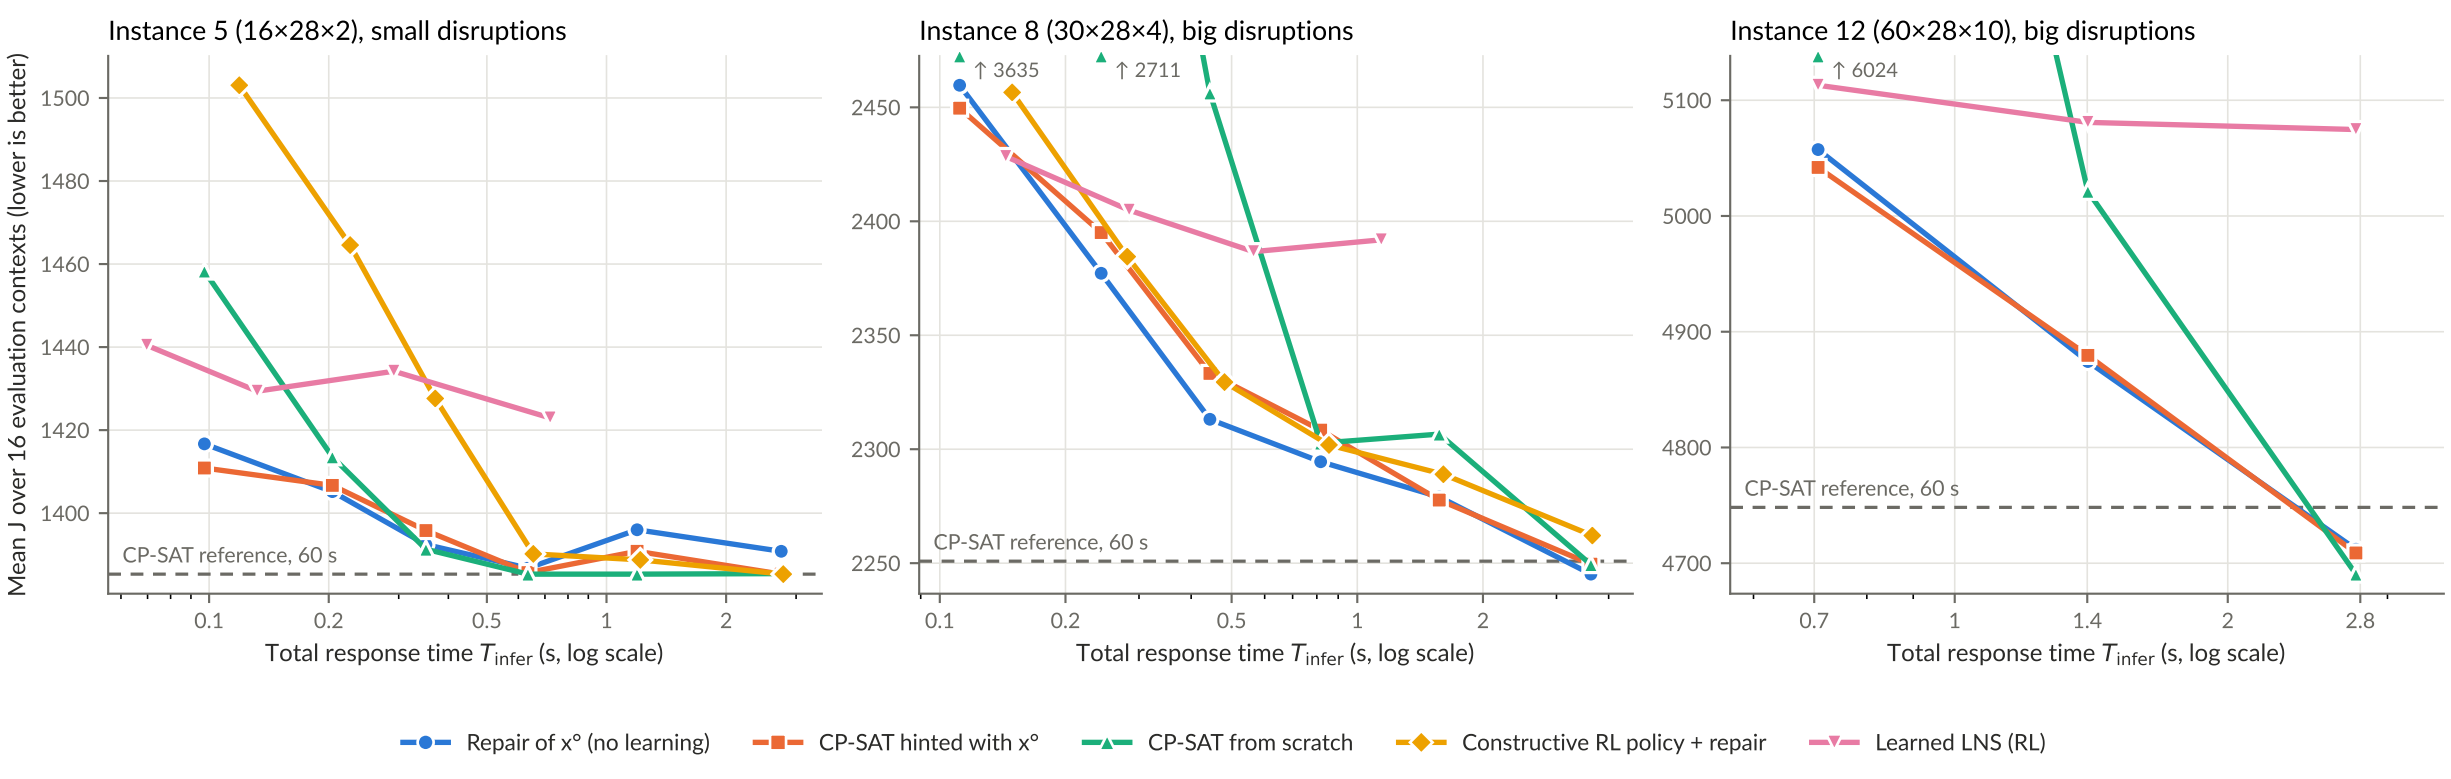
Instance 5 (16×28×2), small disruptions
Instance 8 (30×28×4), big disruptions
Instance 12 (60×28×10), big disruptions
Mean J over 16 evaluation contexts (lower is better)
1500
1480
1460
1440
1420
1400
2450
2400
2350
2300
2250
5100
5000
4900
4800
4700
↑ 3635
↑ 2711
↑ 6024
CP-SAT reference, 60 s
0.1
0.2
0.5
1
2
0.7
1.4
2.8
Total response time $T_{infer}$ (s, log scale)
Repair of x° (no learning)
CP-SAT hinted with x°
CP-SAT from scratch
Constructive RL policy + repair
Learned LNS (RL)